

This site uses cookies to improve the user experience.

Java Concurrency

1. [Java Concurrency / Multithreading Tutorial](#)
2. [Multithreading Benefits](#)
3. [Multithreading Costs](#)
4. [Concurrency Models](#)
5. [Same-threading](#)
6. [Concurrency vs. Parallelism](#)
7. [Creating and Starting Java Threads](#)
8. [Race Conditions and Critical Sections](#)
9. [Thread Safety and Shared Resources](#)
10. [Thread Safety and Immutability](#)
11. [Java Memory Model](#)
12. [Java Synchronized Blocks](#)
13. [Java Volatile Keyword](#)
14. [Java ThreadLocal](#)
15. [Thread Signaling](#)
16. [Deadlock](#)
17. [Deadlock Prevention](#)
18. [Starvation and Fairness](#)
19. [Nested Monitor Lockout](#)
20. [Slipped Conditions](#)
21. [Locks in Java](#)
22. [Read / Write Locks in Java](#)
23. [Reentrance Lockout](#)
24. [Semaphores](#)
25. [Blocking Queues](#)
26. [Thread Pools](#)
27. [Compare and Swap](#)
28. [Anatomy of a Synchronizer](#)
29. [Non-blocking Algorithms](#)
30. [Amdahl's Law](#)
31. [Java Concurrency References](#)

Java Volatile Keyword

- [The Java volatile Visibility Guarantee](#)
 - [Still a Race Condition](#)
- [The Java volatile Happens-Before Guarantee](#)
- [volatile is Not Always Enough](#)
- [When is volatile Enough?](#)
- [Performance Considerations of volatile](#)



Jakob Jenkov
Last update: 2018-02-23



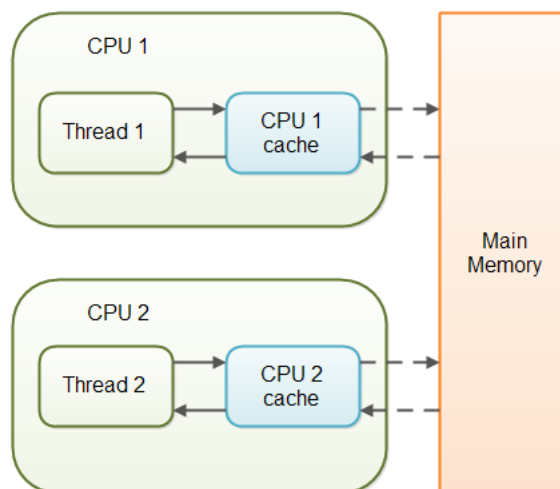
The Java `volatile` keyword is used to mark a Java variable as "being stored in main memory". More precisely that means, that every read of a volatile variable will be read from the computer's main memory and not from the CPU cache, and that every write to a volatile variable will be written to main memory, not just to the CPU cache.

Actually, since Java 5 the `volatile` keyword guarantees more than just that volatile variables are written and read from main memory. I will explain that in the following sections.

The Java volatile Visibility Guarantee

The Java `volatile` keyword guarantees visibility of changes to variables across threads. This may sound a bit abstract, so let me elaborate.

In a multithreaded application where the threads operate on non-volatile variables, each thread may copy variables from main memory into a CPU cache while working on them, for performance reasons. If your computer contains more than one CPU, each thread may run on a different CPU. That means, that each thread may copy the variables into the CPU cache of different CPUs. This is illustrated here:



With non-volatile variables there are no guarantees about when the Java Virtual Machine (JVM) reads data from main memory into CPU caches, or writes data from CPU caches to main memory. This can cause several problems which I will explain in the following sections.

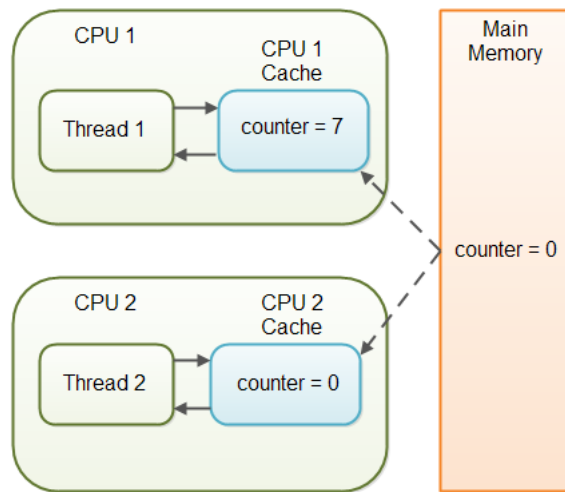
Imagine a situation in which two or more threads have access to a shared object which contains a color variable declared like this:

```
public class SharedObject {
```

```
}
```

Imagine too, that only Thread 1 increments the counter variable, but both Thread 1 and Thread 2 may read the counter variable from time to time.

If the counter variable is not declared `volatile` there is no guarantee about when the value of the counter variable is written from the CPU cache back to main memory. This means, that the counter variable value in the CPU cache may not be the same as in main memory. This situation is illustrated here:



The problem with threads not seeing the latest value of a variable because it has not yet been written back to main memory by another thread, is called a "visibility" problem. The updates of one thread are visible to other threads.

By declaring the counter variable `volatile` all writes to the counter variable will be written back to main memory immediately. Also, all reads of the counter variable will be read directly from main memory. Here is how the `volatile` declaration of the counter variable looks:

```
public class SharedObject {  
    public volatile int counter = 0;  
}
```

Declaring a variable `volatile` thus *guarantees the visibility* for other threads of writes to that variable.

Still a Race Condition

Please notice, that the above example counter is just an example I used to explain how the Java `volatile` keyword affects when values are read and written from and to main memory. However, simply declaring the counter variable `volatile` is not enough to create a fully thread safe shared counter. Let me explain why:

When two threads first read, test and then write a new value based on an older value to a variable, even if the variable is declared `volatile` you will still have a race condition. Imagine this situation where two threads named T1 and T2 are reading, testing and writing to a shared `volatile` variable:

```
T1: Read counter variable - value is 6  
T2: Read counter variable - value is 6  
  
T1: if counter < 100 add 1 to counter  
   -> so adds 1 to counter which is now 7 in T1  
   -> counter is immediately written back to main memory  
  
T2: if counter < 100 add 1 to counter  
   -> so adds 1 to counter which is now 7 in T2  
   -> counter is immediately written back to main memory  
  
Value of counter in main memory is now 7,  
but should have been 8
```

Notice how the two threads T1 and T2 read the shared counter variable directly from main memory, and yet because the read-test-write operation not carried out as a single atomic operation, both threads read the pre-incremented value of the shared counter which is 7. Both threads then increment the counter and write their incremented value directly back to main memory, where T2 does not see T1's incrementation of the counter in main memory. T2 effectively "ignores" T1's incrementation, and sets the value to 7, even though the correct value would

So, as you can see, if two threads are incrementing the shared counter simultaneously, an error can occur, even if both write the incremented value back to main memory as soon as it has been incremented.

To make the read-test-write operation a single atomic operation, you could use one of the atomic primitives in the `java.util.concurrent` package. For instance, you could use [AtomicInteger](#) or [AtomicLong](#).

I also have a discussion of atomic compare-and-swap operations in my [Java compare and swap tutorial](#).

The Java volatile Happens-Before Guarantee

Since Java 5 the `volatile` keyword guarantees more than just the reading from and writing to main memory of variables. Actually, the `volatile` keyword guarantees this:

- If Thread A writes to a volatile variable and Thread B subsequently reads the same volatile variable, then all variables visible to Thread A *before* writing the volatile variable, will also be visible to Thread B *after* it has read the volatile variable.
- The reading and writing instructions of volatile variables cannot be reordered by the JVM (the JVM may reorder instructions for performance reasons as long as the JVM detects no change in program behaviour from the reordering). Instructions before and after can be reordered, but the volatile read or write cannot be mixed with these instructions. Whatever instructions follow a read or write of a volatile variable are guaranteed to happen after the read or write.

These statements require a deeper explanation.

When a thread writes to a volatile variable, then not just the volatile variable itself is written to main memory. Also all other variables changed by the thread before writing to the volatile variable are also flushed to main memory. When a thread reads a volatile variable it will also read all other variables from main memory which were flushed to main memory together with the volatile variable.

Look at this example:

```
Thread A:
    sharedObject.nonVolatile = 123;
    sharedObject.counter     = sharedObject.counter + 1;

Thread B:
    int counter     = sharedObject.counter;
    int nonVolatile = sharedObject.nonVolatile;
```

Since Thread A writes the non-volatile variable `sharedObject.nonVolatile` before writing to the volatile `sharedObject.counter`, then both `sharedObject.nonVolatile` and `sharedObject.counter` are written to main memory when Thread A writes to `sharedObject.counter` (the volatile variable).

Since Thread B starts by reading the volatile `sharedObject.counter`, then both the `sharedObject.counter` and `sharedObject.nonVolatile` are read from main memory into the CPU cache used by Thread B. By the time Thread B reads `sharedObject.nonVolatile` it will see the value written by Thread A.

Developers may use this extended visibility guarantee to optimize the visibility of variables between threads. Instead of declaring each and every variable `volatile`, only one or a few need be declared `volatile`. Here is an example of a simple `Exchanger` class written after that principle:

```
public class Exchanger {

    private Object  object      = null;
    private volatile hasNewObject = false;

    public void put(Object newObject) {
        while(hasNewObject) {
            //wait - do not overwrite existing new object
        }
        object = newObject;
        hasNewObject = true; //volatile write
    }

    public Object take(){
        while(!hasNewObject){ //volatile read
            //wait - don't take old object (or null)
        }
        Object obj = object;
        hasNewObject = false; //volatile write
        return obj;
    }
}
```

Thread A may be putting objects from time to time by calling `put()`. Thread B may take objects from time to time by calling `take()`. This `Exchanger` can work just fine using a `volatile` variable (without the use of synchronized blocks), as long as only Thread A calls `put()` and only Thread B calls `take()`.

However, the JVM may reorder Java instructions to optimize performance, if the JVM can do so without

```

while(hasNewObject) {
    //wait - do not overwrite existing new object
}
hasNewObject = true; //volatile write
object = newObject;

```

Notice the write to the `volatile` variable `hasNewObject` is now executed before the new object is actual set. To the JVM this may look completely valid. The values of the two write instructions do not depend each other.

However, reordering the instruction execution would harm the visibility of the object variable. First of all, Thread B might see `hasNewObject` set to `true` before Thread A has actually written a new value to the object variable. Second, there is now not even a guarantee about when the new value written to `object` will be flushed back to main memory (well - the next time Thread A writes to a volatile variable somewhere...).

To prevent situations like the one described above from occurring, the `volatile` keyword comes with a "*happens before guarantee*". The happens before guarantee guarantees that read and write instructions of volatile variables cannot be reordered. Instructions before and after can be reordered, but the volatile read/write instruction cannot be reordered with any instruction occurring before or after it.

Look at this example:

```

sharedObject.nonVolatile1 = 123;
sharedObject.nonVolatile2 = 456;
sharedObject.nonVolatile3 = 789;

sharedObject.volatile      = true; //a volatile variable

int someValue1 = sharedObject.nonVolatile4;
int someValue2 = sharedObject.nonVolatile5;
int someValue3 = sharedObject.nonVolatile6;

```

The JVM may reorder the first 3 instructions, as long as all of them *happens before* the volatile write instruction (they must all be executed before the volatile write instruction).

Similarly, the JVM may reorder the last 3 instructions as long as the volatile write instruction *happens before* all of them. None of the last 3 instructions can be reordered to before the volatile write instruction.

That is basically the meaning of the Java volatile happens before guarantee.

volatile is Not Always Enough

Even if the `volatile` keyword guarantees that all reads of a volatile variable are read directly from main memory, and all writes to a volatile variable are written directly to main memory, there are still situations where it is not enough to declare a variable `volatile`.


In the situation explained earlier where only Thread 1 writes to the shared `counter` variable, declaring `counter` variable `volatile` is enough to make sure that Thread 2 always sees the latest written value.


In fact, multiple threads could even be writing to a shared `volatile` variable, and still have the correct value stored in main memory, if the new value written to the variable does not depend on its previous value. In other words, if a thread writing a value to the shared `volatile` variable does not first need to its value to figure out its next value.

As soon as a thread needs to first read the value of a `volatile` variable, and based on that value generate a new value for the shared `volatile` variable, a `volatile` variable is no longer enough to guarantee correct visibility. The short time gap in between the reading of the `volatile` variable and the writing of its new value, creates a **race condition** where multiple threads might read the same value of the `volatile` variable, generate a new value for the variable, and when writing the value back to main memory - overwrite each other's values.

The situation where multiple threads are incrementing the same counter is exactly such a situation where a `volatile` variable is not enough. The following sections explain this case in more detail.

Imagine if Thread 1 reads a shared `counter` variable with the value 0 into its CPU cache, increment it 1 and not write the changed value back into main memory. Thread 2 could then read the same `counter` variable from main memory where the value of the variable is still 0, into its own CPU cache. Thread 2 could then also increment the counter to 1, and also not write it back to main memory. This situation is illustrated in the diagram below:





Thread 1 and Thread 2 are now practically out of sync. The real value of the shared counter variable should have been 2, but each of the threads has the value 1 for the variable in their CPU caches, and in main memory the value is still 0. It is a mess! Even if the threads eventually write their value for the shared counter variable back to main memory, the value will be wrong.

When is volatile Enough?

As I have mentioned earlier, if two threads are both reading and writing to a shared variable, then using the `volatile` keyword for that is not enough. You need to use a `synchronized` block in that case to guarantee that the reading and writing of the variable is atomic. Reading or writing a volatile variable does not block threads reading or writing. For this to happen you must use the `synchronized` keyword around critical sections.

As an alternative to a synchronized block you could also use one of the many atomic data types found in the `java.util.concurrent` package. For instance, the `AtomicLong` or `AtomicReference` or one of the others.

In case only one thread reads and writes the value of a volatile variable and other threads only read the variable, then the reading threads are guaranteed to see the latest value written to the volatile variable. Without making the variable volatile, this would not be guaranteed.

The `volatile` keyword is guaranteed to work on 32 bit and 64 bit variables.

Performance Considerations of volatile

Reading and writing of volatile variables causes the variable to be read or written to main memory. Reading from and writing to main memory is more expensive than accessing the CPU cache. Accessing volatile variables also prevents instruction reordering which is a normal performance enhancement technique. Thus, you should only use volatile variables when you really need to enforce visibility of variables.

Next: [Java ThreadLocal](#)

Share

Tweet



Jakob Jenkov

